


Practical Problem Solving ■

From Computational Thinking to Programs
with Python



Practical Problem Solving

From Computational Thinking to Programs with Python

Igor Benav

Contents

Preface	4
About This Book	4
What You'll Build	5
How to Use This Book	5
Prerequisites	6
1 The Foundations of Computational Thinking	7
1.1 Finding a Name in a Phone Book	7
1.2 What Makes an Algorithm?	12
1.3 On Precision	13
1.4 Breaking Problems Apart	14
1.5 Abstraction	14
1.6 From Algorithms to Programs	15
1.7 Why Python?	17
1.8 What Makes a Problem Computational?	18
1.9 Chapter Summary	18
1.10 Exercises	19
2 Getting Started with Python	21
2.1 Installing Python	21
2.2 Two Ways to Run Code	22
2.3 Telling the Computer What to Show	23
2.4 Remembering Values	24
2.5 Asking the User	25
2.6 Numbers and Text	25
2.7 Hands-On: Personal Receipt Printer	26
2.8 Chapter Summary	28
2.9 Exercises	28
3 Types, Conversions, and Formatting	30
3.1 Why the Receipt Crashed	30
3.2 Python's Core Types	31
3.3 The Same Operator, Different Behavior	32
3.4 Why $0.1 + 0.2$ Isn't 0.3	33
3.5 Type Conversion	35
3.6 Formatting Output with F-Strings	36
3.7 Working with Strings	38

3.8	Hands-On: Receipt Printer, Completed	39
3.9	Chapter Summary	43
3.10	Exercises	43
4	Making Decisions	45
4.1	True, False, and Comparisons	45
4.2	The if Statement	46
4.3	Choosing Between Two Paths	47
4.4	Multiple Conditions with elif	48
4.5	Combining Conditions	50
4.6	Nested Conditionals	52
4.7	Hands-On: Receipt Printer with Validation and Discounts	53
4.8	Chapter Summary	57
4.9	Exercises	58

Preface

Beginner programming books usually teach you a language - they walk through the syntax and the features and hope problem-solving follows, but it usually doesn't. You finish knowing what a `for` loop is without quite knowing when to reach for one, able to read code someone else wrote and stuck the moment you face an empty file.

This book starts from problems instead, with Python introduced as the tool to solve them - every concept arrives when a problem demands it rather than when a language tour happens to reach it. The aim is to think clearly enough about a problem that the code follows on its own.

This matters more now than it did a few years ago, because AI tools like ChatGPT can generate working code from a prompt, and that moved the difficulty somewhere else: syntax got cheap, while knowing what to ask for stayed expensive, along with the judgment to tell a good solution from one that stops working the moment the input gets big - and building that judgment is what this book is after. Using AI to help write code is perfectly fine - as long as you can look at what it produced, decide whether it's good or broken, and know what to do about it. The basic rule is: if you don't even know what to ask AI to find problems, you shouldn't be using it.

About This Book

The book is opinionated: we use Python throughout, `uv` for package management, and the REPL for quick experiments. Rather than survey everything Python can do, it commits to one path, from "I've never programmed" to "I can solve real problems with code," so you have solid ground to stand on when you branch out later.

Computer science ideas like algorithms, abstraction, and decomposition show up at the moment they start mattering: you'll understand binary search not because a chapter is titled "Algorithms" but because you're trying to find a name in a phone book and the obvious approach is too slow.

One trade runs through the whole book: you can usually make a program faster by letting it remember more, and make it use less memory by letting it redo work. The early chapters bump into this trade by accident; around the middle of the book you get the standard language for measuring it, and "is this solution any good" turns into a question with an actual answer.

What You'll Build

The book follows a single program: a receipt printer, started in Chapter 2 and still growing in the final chapter. At the end of each chapter it hits a wall - something it genuinely can't do yet - and the next chapter exists because of that wall.

In Chapter 2 you build the first version, which asks for a store name, an item, a quantity, a price, and prints something that looks like a receipt. What it can't do is math: `input()` hands you text, and the total stays uncalculated. Chapter 3 fixes that with types and conversion, and the receipt gets a real total at the bottom, tax included.

But everything is hardcoded at exactly three items, which is what Chapter 4's conditionals start to fix: the receipt reacts to what it's given, skipping a zero-quantity item, discounting a large subtotal. Then Chapter 5 brings loops, and the three-item limit disappears entirely.

By the later chapters the receipt has a product catalog behind it, an undo for mistakes, sales records that survive a restart, and error handling that shrugs off garbage input; near the end it gets reorganized into a packaged project with proper classes. Then the final chapter drops the guided thread: you get a fresh problem, no receipt in sight, and you solve it from scratch the way the whole book taught - understand it first, build the obvious version, measure where it strains, reach for the structures that fix it, and finish knowing what your solution costs. The same habit carries into other languages, and into telling an AI what to build and then judging what comes back.

How to Use This Book

The chapters are meant to be read in order, each one builds on the previous. If you skip types, conditionals won't make sense; skip conditionals and loops will lose you. And this applies even if you already know some Python - the chapters aren't organized around syntax, they're organized around problems, so a chapter you skip isn't just a feature you already know, it's a wall you never hit and a fix you never needed. Skim fast where the code is familiar, don't skip.

Each chapter opens with a problem you can't solve yet and closes once you can. Inside a chapter the order is always motivation first: the problem we're facing, then the Python that addresses it, then a hands-on where it goes into the receipt printer (or a neighboring program) - so you'll know why something exists before you're asked to use it.

Type the code yourself instead of pasting it - the mistakes you make along the way are half the lesson - and mess with things while you're at it: change a number, delete a line, read what Python complains about.

When you get stuck, resist asking AI for the solution - a hint is fine if you need one, but do the thinking yourself, because that's the part you *can't* delegate. Struggling is normal; it's also the difference between actually learning and just feeling like you are. Chapter 2 makes a related recommendation during setup: no AI-assisted editors until you can judge what they produce.

Prerequisites

This book assumes no prior programming experience, and the equipment list is short:

- A computer running Windows, macOS, or Linux
- A code editor (VS Code works well)
- Willingness to experiment and read error messages instead of panicking

Even if you've never written a line of code before, this book will get you there. Chapter 1 starts with what computer science actually is, Chapter 2 gets Python running.

1 The Foundations of Computational Thinking

People assume computer science is about computers the way they assume astronomy is about telescopes. The telescope is the tool, but the stars are the subject. In computer science, the computer is the tool - the subject is problems: how to define them precisely, how to solve them efficiently, and occasionally how to prove they can't be solved at all.

This distinction shapes everything that follows: you could spend a career writing code without doing much computer science, and you could do important computer science with no code at all. What defines the field isn't the medium, it's the questions: What exactly is this problem? Is there an algorithm to solve it? Is there a better one? What does it even mean being better?

Let's start with a problem.

1.1 Finding a Name in a Phone Book

You need to find "Guido" (the creator of Python) in a programmer's paper phone book with thousands of names, you don't have a search box. How do you do it?

Phone books aren't really used anymore, but the problem maps cleanly to any situation where you're searching a list without built-in search: finding a contact in an old phone, looking up a word in a paper dictionary, scanning a sorted spreadsheet.

One approach: start at the first page. Check the first name; is it "Guido"? No. Check the second. No. Keep going until you find it or run out of names.




Figure 1.1: Linear Search Visualization

This works, it will find the name if it's there since you're potentially scanning every element in the phone book. But notice what happens as the phone book grows: a phone book with 1,000 names might require 1,000 checks, a million names: a million checks. If checking one name takes one second, a million-entry city phone book could take nearly two weeks assuming you get no sleep. That's a solution, but not a practical one.

Now think about how you actually use a phone book. You don't start at page one, you flip to roughly the middle. Why is that? We are unconsciously using the fact that the book is alphabetically ordered, and if we're searching for a name that starts with "M" it's probably closer to the center. Let's also use this fact in our method:

1. Open to the middle of the phone book
2. Check the name at that position
3. If it's "Guido," you're done
4. If "Guido" comes before that name alphabetically, search only the first half
5. If "Guido" comes after, search only the second half
6. Repeat with the chosen half, again splitting it in the middle
7. Continue until you find the name or determine it's not there

Let's trace this on a 20-name phone book. Open to the middle and you land on "James."



Figure 1.2: Binary Search Step 1

“Guido” comes before “James” alphabetically. Discard the second half. Ten names left. Open to the middle of those: “Donald.”




Figure 1.3: Binary Search Step 2

“Guido” comes after “Donald.” Down to 5 names, all between “Donald” and “James.” Middle of that section: “Frances.”



Figure 1.4: Binary Search Step 3

“Guido” comes after “Frances.” Two or three names left. Check the middle: “Guido.” Found it.



Figure 1.5: Binary Search Step 4

Four checks. The search space halved at every step: 20 names, then 10, then 5, then 2–3, then done. Let’s compare the two approaches across phone books of different sizes:

Phone Book Size	Linear Search	Binary Search
1,000 names	1,000 checks	10 checks
1,000,000 names	1,000,000 checks	20 checks
8,000,000 names	8,000,000 checks	23 checks

Eight times as many names adds millions of checks to linear search, but only three more steps to binary search. That gap only widens with size.

Figure 1.6: Search Algorithms Comparison

For 8 million names, if each check takes one second: the first approach takes at most 8 million seconds (over three months), the second takes at most 23 seconds. Both approaches are correct, but one of them is actually usable.

Notice what “better” actually meant there: not that binary search wins on one particular phone book - it’s that as the book grows, binary search’s cost barely grows with it, while linear search’s climbs right alongside the size. That’s almost always what we mean when we call one algorithm better than another: not that it’s faster on some specific input, but that it’s slower to get worse as the input grows. It’s a small shift in the question - from “how long does this take?” to “how does the cost grow as the problem grows?”. We’ll come back to it later with a language of its own. Steps aren’t the only thing a solution spends, either - memory is the other one. A recurring move in everything ahead is trading one for the other: spending some memory to make an algorithm faster, or accepting more steps to use less space.

Each time you double the phone book’s size, the first approach takes twice as long, the second only needs one more step. This is what computer scientists call *logarithmic scaling*, one of the most powerful ideas in algorithm design.

We’ve been talking about “approaches” and “methods.” Computer science has a precise word for what we’ve described: an **algorithm**.

1.2 What Makes an Algorithm?

Both phone book approaches share certain properties: they’re sequences of steps, they take an input (a name to find, a phone book to search), they produce an output (the name’s location, or the fact that it’s

not there), they terminate: you always finish, whether or not the name exists.

That's what an **algorithm** is: a finite, well-defined sequence of computational steps that takes an input and produces an output. The word "finite" is important here: an algorithm must terminate. "Well-defined" is also important: each step must be unambiguous: precise enough that a machine (or a person following instructions mechanically) could execute it without guessing.

The phone book showed us two algorithms for the same problem - the first is called **linear search**: check each element in sequence, the second is called **binary search**: use the sorted structure to halve the search space at every step. Both are correct but one is dramatically better for our use-case. For small inputs, efficiency barely matters, but for large inputs, algorithm choice determines whether a solution is practical at all.

Let's think about the two questions we just asked about these algorithms. First: is it correct - does it actually produce the right answer, every time, not just on the examples we happened to try? Both of our searches pass this one. Second: how does it scale - what happens to its cost as the input grows? Here the two split apart completely. Almost everything you'll do from here on is some version of these two questions: get it right, then get it to scale.

Binary search also reveals something subtler: it works because the phone book is sorted. If names were in random order, you couldn't discard half the book with each comparison, so the structure of the data enabled a shortcut. You'll see this pattern showing up everywhere: finding structure in your problem often unlocks better algorithms. The algorithm and the way the data is arranged aren't separate concerns - they're two halves of the same thing. The right structure is often what makes the algorithm you want possible in the first place. A big part of what's ahead is really about organizing data so that the good algorithms open up to you - lists, dictionaries, sets, and the rest make better solutions possible. That structure wasn't free, though - someone sorted the book once so that every search afterward could be cheap. Paying a cost up front to make repeated work fast is a trade you'll see constantly.

But understanding an algorithm and being able to explain it precisely are two different things. You followed binary search intuitively, but could you write it down well enough for someone who has never seen a phone book?

1.3 On Precision

Consider an algorithm for making a sandwich:

1. Take two slices of bread
2. Spread mayonnaise on one slice
3. Add lettuce, tomato, and cheese
4. Place the second slice on top
5. Cut the sandwich in half

A human understands this instantly, but read it again as if you've never seen a sandwich. Which side of the bread gets the mayonnaise? How much? How should the tomato be sliced? Where does each ingredient go? What does "cut in half" mean: diagonally? Down the middle?

You fill in these gaps without even thinking; you've made sandwiches, you know what bread is, you have common sense - a computer has none of that. It needs exact quantities, exact positions, exact actions. Every unstated assumption, every common-sense inference, has to be made explicit.

Go back to binary search. "Open to the middle of the phone book" is clear enough for you. But a computer needs to know: how do you compute "the middle" of a list? If the list has an even number of entries, which of the two middle entries do you check? When you "discard the second half," do you include the entry you checked, or not? These details don't change whether the algorithm works in principle, but they're the difference between a rough spec and an actual implementation, you'll need to deal with these things.

This is what makes programming difficult when you're starting out, you're shifting from implicit to explicit: you have to say what you mean completely, not rely on the reader to fill in the rest. A machine that follows instructions exactly (with no assumptions, no improvisation) can do it billions of times without fatigue. But it will only do exactly what you said, not what you meant.

Our search problem had a really simple scope, but what happens when the problem is bigger than a single search?

1.4 Breaking Problems Apart

Finding a name in a phone book is one task, building a contacts app involves many: adding new contacts, sorting them, searching by name, searching by number, handling duplicates, displaying results. Each of these is its own problem with its own algorithm. The skill is recognizing where a large problem divides into smaller, independent pieces.

This is called **decomposition**: breaking a complex problem into subproblems, each manageable on its own. It mirrors what binary search does to the phone book - split the problem in half, solve the smaller version, a strategy common enough to have a name: divide and conquer - but applied to the problem's structure rather than its data.

Decomposition alone isn't enough, though. Once you've broken a problem into pieces, you need a way to work on each piece without drowning in the details of all the others.

1.5 Abstraction

When you described binary search to a friend, you probably said something like "open to the middle, check the name, throw away half." You didn't explain how pages work or how alphabetical ordering is defined. You worked at a level where those details didn't really matter to understand the whole.

This is **abstraction**: hiding what you don't need to think about so you can focus on what you do.

When you use a smartphone, you don't think about transistors or voltage levels or memory addresses, you simply tap an icon. The icon is an abstraction: an interface that hides enormous complexity underneath. You don't need to understand what's happening at the hardware level to use the app, the abstraction lets you work at a higher level.

Programming languages work the same way. When you call `print()` in Python, you're not writing the machine instructions that move characters to a screen buffer, you're using an abstraction that handles all of that for you. Python itself is an abstraction over assembly code, which is an abstraction over machine code, which is an abstraction over electrical signals running through the processor.

Binary search has layers too. At the top: "halve the search space until you find the target". One level down: "compare the target to the middle element and decide which half to keep". One level further: "compute the index of the middle element given the current bounds". You can think about the strategy without thinking about index arithmetic, and you can think about index arithmetic without thinking about how memory stores the list. Each layer trusts the one below it.

This layering is how complex systems can get built at all: you can work at one level without understanding (or even having to think about) the levels below. When you need to go deeper, you can, but most of the time, the abstraction holds and you don't have to.

We've been describing algorithms in words and diagrams, but at some point they need to run on a computer. That requires translating them into a language a computer can execute.

1.6 From Algorithms to Programs

An **algorithm** is a logical sequence of steps to solve a problem: language-independent, the idea itself; a **program** is that algorithm implemented in a specific programming language: one particular expression of the idea.

Binary search is an algorithm, and you could implement it in Python, Java, C++, or any other language. The algorithm stays the same; the syntax changes. Programming is the act of transcribing an algorithm into a language a computer can execute.

This wasn't always done in high-level languages though, early programmers gave instructions directly in machine code: the ones and zeros the processor actually understands. A program to display "Hello, World!" in machine code looks like this:

```
B4 03 CD 10 B0 01 B3 0A B9 0D 00 BD 13 01 B4 13 CD
10 C3 48 65 6C 6C 6F 2C 20 57 6F 72 6C 64 21 0D 0A
```

Broken down, those hexadecimal codes represent specific processor instructions (This is real-mode DOS assembly, from back when programmers worked this close to the metal - it's here to look at, not to run. It won't execute on a modern 64-bit machine as-is):


```

B4 03      ; MOV AH, 03h      (Get cursor position)
CD 10      ; INT 10h          (BIOS video interrupt)
B0 01      ; MOV AL, 01h      (Write mode parameter)
B3 0A      ; MOV BL, 0Ah      (Color attribute)
B9 0D 00   ; MOV CX, 000Dh    (Character count - 13 for "Hello, World!")
BD 13 01   ; MOV BP, 0113h    (Offset to string data)
B4 13      ; MOV AH, 13h      (Write string function)
CD 10      ; INT 10h          (BIOS video interrupt)
C3         ; RET              (Return)

```

And the text data itself, encoded as numbers:

```
48 65 6C 6C 6F 2C 20 57 6F 72 6C 64 21 0D 0A
```

You don't need to understand this, you need to feel the gap between that and the Python equivalent:

```
print("Hello, World!")
```

Same result, one line versus dozens of cryptic instructions. That gap is what high-level programming languages exist to close (and it's abstraction at work: `print()` hides all of that machine-level complexity behind a single word).

Here's another example, this algorithm finds the maximum value in a list:

1. Start with the first number as the current maximum
2. For each remaining number, if it's larger than the current maximum, it becomes the new maximum
3. After checking all numbers, the current maximum is the answer

And here's what it looks like in Python (don't worry about the syntax for now):

```

def find_maximum(numbers):
    if not numbers:
        return None

    maximum = numbers[0]

    for number in numbers[1:]:
        if number > maximum:
            maximum = number

    return maximum

my_numbers = [42, 17, 8, 94, 23, 61]
print(f"The maximum value is {find_maximum(my_numbers)}")

```

The algorithm is the logic, the Python code is one way to express it. Keeping these separate in your head is important because when your program gives the wrong answer, the problem is either in the algorithm

(the logic is wrong) or in the implementation (the code doesn't correctly express the logic) - different bugs with different fixes.

The choice of which language to use is a separate question, and for this book, it's already been made for you.

1.7 Why Python?

Python was created by Guido van Rossum in the late 1980s. Its central design philosophy, summarized in a document called "The Zen of Python," includes lines like:

- Beautiful is better than ugly
- Explicit is better than implicit
- Simple is better than complex
- Readability counts

These rules make Python code read a lot like pseudocode: instructions close enough to plain English that the algorithm is visible in the implementation. For learning, that's valuable; you can focus on the problem without fighting the language.

Python is also practical since it's the dominant language in data science, machine learning, scientific computing, and automation. Learning it doesn't only teach you to program; it gives you access to a vast ecosystem of tools built by other people solving related problems.

It has real limitations though, and you should know them to be able to decide when it's the best option. Python is interpreted: instead of being compiled all the way to machine code ahead of time, your code is translated into an intermediate form (called bytecode) and executed step by step as it runs. That extra work at runtime is a big part of why compiled languages like C++ or Rust are typically much faster. For most applications this doesn't matter, but for performance-critical or low-resource systems, it does. Python also has a mechanism called the Global Interpreter Lock (the **GIL**) that allows only one thread to execute Python code at a time, which limits certain kinds of parallel processing. Note that the GIL can be disabled since python 3.13, but it will take some time until packages adapt and fix the bugs that the GIL was previously suppressing. And some domains (mobile apps, game engines, low-level systems programming) are better served by other languages. Every language is a tradeoff, python is optimized for programmer productivity and the scientific and data ecosystem.

Professional programmers typically know several languages and choose based on the problem. The concepts you'll learn here (algorithms, data structures, abstraction) apply in any language, python is a clear place to start.

We know that the concepts matter more than the tool, but this raises one more question: what kinds of problems can these concepts actually solve?

1.8 What Makes a Problem Computational?

Not every problem can be solved with an algorithm. We've been working with "find a name in a phone book," which is clearly solvable: there's a finite list, a defined target, and a procedure that terminates, but consider "what is the most beautiful painting?" - there's no unambiguous definition of "beautiful," no finite set of candidates everyone agrees on, and no procedure that would produce an answer everyone accepts. It's not a computational problem.

For a problem to be computational, it needs to be well-defined (the problem statement is clear and unambiguous), solvable (there exists a finite sequence of steps that reaches an answer), and finite (the solution must terminate). "What is the sum of all integers from 1 to 100?" is computational. "What career would make someone happiest?" is not.

Computational thinking is the practice of reformulating problems so they become solvable algorithmically. This often involves ignoring irrelevant details (abstraction), breaking the problem into pieces (decomposition), and recognizing when a new problem resembles one you've already solved (pattern recognition). It's the collection of concepts we've been defining throughout this chapter, applied deliberately.

The phone book showed us all of this in miniature: we started with a problem, we found two algorithms and discovered that the choice between them matters enormously, then we saw that the data's structure (sorted order) enabled the better algorithm. We also used decomposition and abstraction to manage complexity, then distinguished the algorithm (the idea) from the program (its implementation in a specific language).

These are the ideas the rest of the book applies. In the next chapter, we move from theory to practice: setting up Python and writing actual programs, but always starting with the thinking.

1.9 Chapter Summary

Computer science is primarily about solving problems, not about computers. The phone book made the stakes concrete: linear search and binary search both find the name, but at scale the difference between them is the difference between seconds and months; algorithm choice dominates.

An algorithm is a finite, well-defined sequence of steps that takes an input and produces an output. "Finite" and "well-defined" are both essential. Binary search also showed that the structure of the data (sorted order) enabled the better algorithm. Finding structure in your problem is often what unlocks a better solution.

The sandwich example exposed the precision gap: humans tolerate vague instructions, computers don't. Programming requires you to close that gap, stating every step explicitly enough that a machine can follow without guessing. Decomposition breaks large problems into manageable pieces. Abstraction lets you work at one level without drowning in the details of the levels below. Both are how complex systems get built at all.

Algorithms are language-independent descriptions of solutions. Programs are their implementation in a specific language. Keeping those two things distinct makes debugging much cleaner: when something goes wrong, is the logic flawed or did the code fail to express the logic correctly?

In the next chapter, we set up Python and write actual programs. The exercises below are worth doing even if they seem obvious; they're not testing what you know, but wiring the circuits.

1.10 Exercises

1. Write a detailed algorithm for a common daily task: making breakfast, getting ready in the morning, or anything you do regularly. Be precise enough that someone who has never done it could follow your instructions without guessing. Once you're done, go back and ask: where did you make assumptions? What did you leave implicit that a computer would need stated explicitly? The gap between your first draft and a truly unambiguous version is the gap that programming requires you to close.
2. Using the sorted list of 16 programming languages below, pick any language other than Python and trace both algorithms to find it. For linear search, check each language in order and count the comparisons; For binary search, use divide-and-conquer and count the comparisons, then write out every step of both. The difference in step counts is the point, but try to also articulate *why* binary search needs so many fewer steps.

1. Ada	5. Cobol	9. JavaScript	13. Python
2. C	6. Fortran	10. Kotlin	14. Ruby
3. C#	7. Go	11. Perl	15. Rust
4. C++	8. Java	12. PHP	16. Swift

3. Choose one of these complex tasks and break it into subtasks, each precise enough to be its own algorithm: planning a vacation, organizing a birthday party, or writing a research paper. The goal is identifying the natural seams in the problem, the places where a complex task divides into independent pieces. How deep do you need to go before each piece feels manageable?
4. For each of the following, decide whether it's computational (solvable by algorithm) or not, and explain your reasoning: (a) finding the shortest route between two cities, (b) determining the most beautiful musical composition, (c) calculating the average temperature for a city over a year, (d) converting currencies based on current exchange rates, (e) deciding which career would make someone happiest. Some of these are less obvious than they first appear.
5. Below is a 28-step grocery shopping trip written out in full. Group the steps into meaningful functions and rewrite the whole trip using only those function names, so the high-level description fits in a few lines; the complexity lives inside the named abstractions, not at the top level. Then ask yourself: does the abstraction actually hide the right things? Are there groupings that feel arbitrary, and why?

1. Write down needed items
 2. Check what you already have at home
 3. Update shopping list
 4. Get reusable bags
 5. Find car keys
 6. Drive to grocery store
 7. Park the car
 8. Grab a shopping cart
 9. Enter the store
 10. Check first item on list
 11. Find the correct aisle
 12. Locate the item
 13. Check price and quality
 14. Place item in cart
 15. Cross item off list
 16. Repeat steps 10-15 until list is complete
 17. Proceed to checkout
 18. Wait in line
 19. Place items on conveyor belt
 20. Pay for groceries
 21. Bag the groceries
 22. Return cart
 23. Walk to car
 24. Load bags into car
 25. Drive home
 26. Carry bags inside
 27. Unpack groceries
 28. Store items in proper places
6. The binary search algorithm in this chapter relies on the phone book being sorted. What would happen if it weren't? Try tracing binary search on an unsorted list of 8 names (make one up). At what point does the algorithm break down? What does this tell you about the relationship between an algorithm and the assumptions it makes about its input?

2 Getting Started with Python

Hopefully you understand binary search and could explain it to someone now, but explaining it to a computer requires a level of precision that English doesn't demand: exact syntax, explicit instructions, no unstated assumptions. In this chapter we'll move from the abstract specification to the concrete implementation.

We'll set up Python, learn how to run code, and build a small program that takes input from a user, processes it, and produces formatted output, which is the structure of most programs you'll ever write - an input-process-output loop:




Figure 2.1: The Input-Process-Output Loop

By the end of the chapter you'll have written one.

2.1 Installing Python

We'll use a single tool for everything: **uv**. It installs Python for you, manages your projects, and keeps each project's dependencies isolated - jobs that traditionally took three or four separate tools. Learning one thing now saves you a lot of confusion later.

Install uv first. It runs on Windows, macOS, and Linux.

On **Windows**, open PowerShell and run:

```
powershell -ExecutionPolicy ByPass -c "irm https://astral.sh/uv/install.ps1 | iex"
```

On **macOS or Linux**, open a terminal and run:

```
curl -LsSf https://astral.sh/uv/install.sh | sh
```

Close and reopen your terminal so it picks up the new command, then check it worked:

```
uv --version
```

If you see a version number, you're set. Notice what you *didn't* have to do: you never installed Python directly. uv does that on demand - the first time a project needs a particular Python version, it fetches it for you. You can also do it explicitly:

```
uv python install
```

For a code editor, Visual Studio Code is the right choice for this book. It's free, runs on all platforms, and has solid Python support through its Python extension. Download it from code.visualstudio.com and install the Python extension from the Extensions panel. If you can't install software locally, pydantic.run is an online Python environment that works in a browser with no setup at all - a good fallback for getting started, though you'll want a local setup for the later chapters.

One recommendation: don't reach for an AI-assisted editor like Cursor or Copilot yet. You'll use these tools eventually - they're genuinely useful - but right now they'll write the code for you before you understand what you're writing, which makes the early chapters pointless. The struggle is a big part of how the ideas actually stick. Later, the right use of these tools is to help you write code faster, not to think for you - and this book is about learning to think. Use a plain editor until you can judge for yourself whether what the AI produces is any good.

Once uv is installed, there are two main ways to run Python code.

2.2 Two Ways to Run Code

The first way is interactive. Open a terminal, run `uv run python`, and you'll see a prompt that looks like `>>>`. This is the **REPL** (Read-Evaluate-Print Loop): you type an expression, Python evaluates it, prints the result, and waits for the next input. (`uv run` just makes sure you're using the Python that uv manages, rather than whatever else might happen to be on your system.)

```
>>> 2 + 3
5
>>> "Hello" + " " + "World"
```

```
'Hello World'
```

The REPL is useful for testing small ideas quickly: does this expression do what I think? You'll use it throughout the book to check your understanding before committing something to a file. Type `exit()` or press Ctrl+D (Ctrl+Z on Windows) to leave it.

The second way is script mode, you write code in a file with a `.py` extension and run the whole file at once. Create a file called `first.py` with this content:

```
2 + 3
"Hello" + " " + "World"
```

Run it with `uv run first.py`. Nothing appears. That's not a bug: in script mode, Python evaluates expressions but doesn't print results unless you explicitly ask it to. This is the first taste of the precision gap from Chapter 1: you probably assumed Python would show you the result. Python did exactly what you said, which was "compute this value", you never said "show it to me."

That's what `print()` is for.

2.3 Telling the Computer What to Show

Fix the script:

```
print(2 + 3)
print("Hello" + " " + "World")
```

Now it outputs:

```
5
Hello World
```

`print()` takes whatever you give it and displays it on screen.

You can print multiple values in a single call by separating them with commas and Python inserts a space between them:

```
print("The answer is", 42)
```

```
The answer is 42
```

By default `print()` adds a newline after each call, you can change that with the `end` parameter: `print("Hello", end=" ")` leaves the cursor on the same line and just adds a space so the next `print()` continues there. The default value of `end` is `"\\n"`, which Python interprets as a new line.

Now we can display things, but everything so far is hardcoded: the program does the same thing every time. To build anything useful, we need to work with values that change, so we need the computer to remember things.

2.4 Remembering Values

A **variable** is a name that refers to a value. You create one with the assignment operator `=`:

```
message = "Hello, Python!"
number = 42
print(message)
print(number)
```

```
Hello, Python!
42
```

The left side is the name, the right side is the value. Once assigned, the name can be used anywhere you'd use the value directly. This is how programs hold onto information: give it a name that references where the value is in the memory, then use that name later.

Variables can be reassigned:

```
count = 10
count = count + 1
print(count)
```

```
11
```

The second line reads the current value of `count`, adds 1, and assigns the result back to `count`. This pattern is constant in programming: it's how counters, accumulators, and running totals work.

Variable names can contain letters, numbers, and underscores, but can't start with a number and can't be Python keywords like `if` or `for`. Python convention is lowercase words separated by underscores: `user_name`, `total_score`, `highest_temperature`. Names are case-sensitive: `score` and `Score` are different variables. Also, use descriptive names (`user_age` instead of `x`); you'll thank yourself later.

Now we can display things and remember them, but the program still does the same thing every time it runs, because every value is written into the code. A useful program responds to whoever is using it, it needs to ask questions.

2.5 Asking the User

`input()` pauses the program, displays a prompt, waits for the user to type something and press Enter, then returns what they typed as a string (how Python stores text).

```
name = input("What is your name? ")
print("Hello, " + name + "!")
```

When this runs, the user sees `What is your name?` and can type a response; whatever they type gets stored in `name` and used in the greeting. This is the first program we've written that behaves differently depending on who runs it.

One thing to keep in mind: `input()` always returns a string, even if the user types a number. If you ask someone for their age and they type 25, you get the string (that is, the text) `"25"`, not the integer number 25. This distinction between text and numbers will be important very soon.

We now have four tools: `print()` to display output, expressions to compute values, variables to remember them, and `input()` to get information from the user. Before we combine them, we need to understand how Python handles the two main kinds of data we'll work with: numbers and text.

2.6 Numbers and Text

Python supports the arithmetic operators you'd expect: `+` for addition, `-` for subtraction, `*` for multiplication, `/` for division. Division always returns a decimal number (one with a whole number part and a fractional part, separated by a decimal point), even when the result is a whole number: `10 / 2` gives `5.0`, not `5`. Two additional operators are worth knowing early: `**` for exponentiation (`2 ** 8` gives 256) and `%` for the remainder after division (`10 % 3` gives 1).

The order of operations follows standard math: exponentiation first, then multiplication and division, then addition and subtraction; parentheses override the default order: `2 + 3 * 4` gives 14, but `(2 + 3) * 4` gives 20 (PEMDAS).

Text in Python is called a **string**, you write strings in either single (`'`) or double (`"`) quotes; the choice doesn't matter as long as you're consistent within a single string. Strings support two operators that look familiar but behave differently. `+` concatenates them: `"Hello" + " " + "World"` gives `"Hello World"`. `*` repeats them: `"ha" * 3` gives `"hahaha"`.

Notice the difference: when Python sees `2 + 3`, it adds the numbers and gets 5. When it sees `"Hello" + " World"`, it glues the strings together and gets `"Hello World"`. The same operator behaves differently depending on the data type, and this is important: `"2" + "3"` gives `"23"`, not 5. The quotes make them strings, so `+` concatenates instead of adding.

Remember that `input()` always returns a string? If the user types 5 and you store it in a variable called `price`, you have the string `"5"`, not the number 5. You can concatenate it with other strings (`"Price: $" + price` gives `"Price: $5"`), but you can't do math with it, at least not yet. We'll fix this in Chapter 3 when we introduce type conversion.

That's the full toolkit for this chapter: output, expressions, variables, input, and the beginnings of understanding data types. Time to use them together.

2.7 Hands-On: Personal Receipt Printer

We'll build a program that collects information about a purchase, formats it, and prints a receipt. It uses every concept from the chapter and produces something you can look at and modify. It also runs straight into a wall that will make the need for Chapter 3 concrete.

One new piece of syntax shows up in the code below: a line that starts with `#` is a **comment**. Python ignores it completely - it's not an instruction, it's a note to yourself (or to whoever reads the code later). We'll use them to label things and leave reminders. Anything after the `#` on that line is invisible to Python.

Set up your project:

```
uv init chapter2
cd chapter2
```

`uv init` creates a small project scaffold - a folder with a couple of config files `uv` uses to track your project. You don't need to worry about what's in them yet; we'll look properly in the modules chapter, near the end of the book. Create a file called `receipt.py` inside this folder. We'll build it up in pieces.

Start by collecting the input:

```
# receipt.py
print("=== Receipt Printer ===", end="\n\n")

store_name = input("Store name: ")
item_name = input("Item name: ")
quantity = input("Quantity: ")
price = input("Price per item: ")
customer_name = input("Customer name: ")
```

If you run it with `uv run receipt.py`, you should be prompted for each piece of information in order. The program collects it but doesn't do anything with it yet.

Now add the receipt output below the input section:

```
# receipt.py
# ...existing code above

print()
print("=" * 30)
print("      " + store_name)
print("=" * 30)
print()
print("Customer: " + customer_name)
```

```

print()
print(item_name)
print(" " + quantity + " x $" + price)
print()
print("-" * 30)
print("  Total items: " + quantity)
print("=" * 30)
print("    Thank you for your purchase!")
print("=" * 30)

```

Run it again and enter some values. If you enter “Bookshop” as the store, “Book” as the item, “2” as the quantity, “14.99” as the price, and “Hannah” as the customer, the receipt looks like:

```

=====
                        Bookshop
=====

Customer: Hannah

Book
  2 x $14.99

-----
Total items: 2
=====
    Thank you for your purchase!
=====

```

Notice that `"=" * 30` produces a row of thirty equals signs. The `*` operator on strings repeats the string, which is a common way to draw lines in terminal output.

Now try to add a total. You have the quantity and the price, the total should be quantity times price. Try adding this line:

```

print("  Total: $" + quantity * price)

```

Run it and you’ll see that Python crashes. The error says something about not being able to multiply a string by a string. `quantity` and `price` are strings (because `input()` always returns strings), and Python doesn’t know how to multiply “2” by “14.99”. You can’t compute “2” * “14.99” any more than you can compute “hello” * “world”, the operation doesn’t make sense for text.

The receipt captures the pattern (collect input, build strings, format output), but it can’t do the one thing a real receipt needs, which is calculating the total. We’re treating numbers as text, and text can’t do arithmetic. Chapter 3 fixes this by introducing data types properly, once you can write `float(price)` * `int(quantity)`, the receipt works for real.

Remove the broken total line for now. The receipt is complete as a string-formatting exercise. Try modifying it: change the formatting, add a date line, widen the receipt; breaking things and seeing what happens is how you actually learn what the rules are.

2.8 Chapter Summary

Setting up Python is overhead, but it only happens from time to time. With `uv` you install one tool and it takes care of fetching Python, creating projects, and keeping them isolated. The two modes of running code (interactive and script) serve different purposes: the REPL for quick experiments, script files for anything you want to keep. The first surprise was that script mode doesn't show results unless you ask: Python does what you say, not what you assume.

`print()` is how programs communicate output; `input()` is how they accept it. Variables store values so they can be reused and combined. These are the minimum tools for the input-process-output loop that most programs follow. The receipt printer demonstrated that, but it also demonstrated a wall: `input()` returns strings, and strings can't do math. The program can display prices but can't add them up, can show quantities but can't multiply them.

That wall is the data type system. In the next chapter, we'll introduce types properly: integers, floats, strings, and the conversions between them.

2.9 Exercises

1. Extend the receipt printer to support three line items instead of one. Ask for a name, quantity, and price for each item separately. Format them as three rows in the receipt body. You'll have three `item_name`, `quantity`, and `price` variables (name them clearly: `item1_name`, `item2_name`, etc.). The total line still can't be calculated yet: display the individual quantities. Which part of the output is hardest to format correctly, and why?
2. Write a program that asks for a sender's name, a recipient's name, and a one-sentence message, then formats and prints a postcard. The postcard should have a border made of a repeated character, the recipient prominently at the top, the message in the middle, and "From: [sender]" at the bottom. The border width should be determined by the length of the message (`"=" * len(message)` gives you a line exactly as wide as the message). What happens if the message is very short or very long?
3. Write a program that asks for someone's first name, last name, and preferred title (Mr., Ms., Dr., etc.), then prints their name formatted three ways: full formal (Dr. Jane Smith), last name first (Smith, Jane), and initials (J.S.). The initials require you to access individual characters in a string: `name[0]` gives you the first character. What happens if you use `name[0]` on an empty string?

4. The `print()` function accepts a `sep` parameter that controls what gets placed between multiple values instead of the default space: `print("a", "b", "c", sep="-")` prints `a-b-c`. Write a program that asks for five words and prints them in three formats: space-separated (default), comma-separated, and with a pipe character between them (`word1|word2|word3`). Now try passing all five variables to a single `print()` call with the appropriate `sep`. How does that compare to building the string manually with `+`?
5. Write a program that generates a name tag for a conference badge. Ask for the person's name, their job title, and the company they work for. The badge should have a fixed width of 30 characters, with each line of text centered. Python strings have a `.center(width)` method: `"hello".center(30)` pads the string with spaces so it's centered in a field of 30 characters. Use this to format the badge. What happens when a name is longer than 30 characters?
6. Write a program that asks for two words and prints every combination of them with a separator between them: the first word, then the second, then both joined, then the second first and the first second, then the first repeated twice, then the second repeated twice. Six lines of output from two inputs. Before writing any code, write out what the six lines should look like for inputs `"black"` and `"bird"`. Does the output match what you expected? The goal isn't to memorize string operations; it's to notice whether your mental model of what `+` and `*` do on strings matches what Python actually does.

3 Types, Conversions, and Formatting

Chapter 2 ended with a crash: the receipt printer collected a quantity and a price from the user, but when you tried to compute the total with `quantity * price`, Python refused - the error said something about not being able to multiply strings. Both values looked like numbers, but `input()` returned them as text, and text can't do arithmetic.

This chapter explains why that happened and fixes it. Along the way, we'll learn how Python distinguishes between different kinds of data, how to convert between them, and how to format output so it looks the way you want. By the end, you'll rebuild the receipt printer with a working total, real currency formatting, and none of the limitations that stopped you last time.

3.1 Why the Receipt Crashed

Let's look at what Python actually told you. When you added the total line and ran the program, you got something like this:

```
Traceback (most recent call last):
  File "receipt.py", line 25, in <module>
    print("  Total: $" + quantity * price)
                        ~~~~~^~~~~~
TypeError: can't multiply sequence by non-int of type 'str'
```

This is a **traceback**: Python's report of where and why your program died. Read it bottom-up; the last line is the diagnosis, the lines above it show the location, and the little `~~~^~~~` markers point at the exact expression that failed. `TypeError` means an operation was applied to the wrong kind of data. "Sequence" is Python's family name for ordered collections of things - a string is a sequence of characters. And "non-int of type 'str'" means the thing you tried to multiply by wasn't an integer, it was another string.

Decoded: Python knows how to multiply a string by an integer (that's the repetition you used for `"=" * 30`), but multiplying a string by a string has no defined meaning. You expected `quantity * price` to multiply 2 by 14.99, but `input()` always returns strings, so `quantity` held "2" and `price` held "14.99" - and `"2" * "14.99"` makes no more sense than `"hello" * "world"`.

This isn't a bug in Python, Python treats "2" and 2 as fundamentally different things: one is text (a sequence of characters that happens to look like a number), the other is a number (a value you can do

arithmetic with). They're stored differently, they support different operations, and Python won't silently guess which one you meant.

The category of data a value belongs to is called its **type**, and understanding types is how you predict what Python will do with your data - and why it sometimes refuses. Learning to read tracebacks like the one above is part of the same skill; you'll see many more of them, and the last line almost always tells you what went wrong.

3.2 Python's Core Types

If you open the REPL (`uv run python`) and try this, you'll get these results:

```
>>> type(42)
<class 'int'>
>>> type(3.14)
<class 'float'>
>>> type("hello")
<class 'str'>
>>> type(True)
<class 'bool'>
```

The `type()` function tells you what kind of data you're working with. Python has four types you'll use constantly:

An **integer** (`int`) is a whole number: 42, -7, 0, 1000000; no decimal point. Python handles integers of any size; you'll never overflow (try to store a value larger than that chunk of memory can hold), no matter how large the number gets.

A **float** is a number with a decimal point: 3.14, -0.001, 1.0. The name comes from "floating-point," a reference to how the decimal point can appear at different positions. Even `1.0` is a float, not an integer - having the decimal point is what makes the difference.

A **string** (`str`) is the text type you met in Chapter 2: any sequence of characters enclosed in quotes. `"42"` is a string, `42` is an integer, and the quotes are what visibly separate them.

A **boolean** (`bool`) has only two possible values: `True` and `False`. You'll use these heavily once we introduce conditionals in Chapter 4; for now, know they exist.

There's also `None`, a special value that represents "nothing" or "no value" - Python's way of saying a variable exists but doesn't have meaningful data yet.

The type of a value determines what you can do with it, and the same operator can behave completely differently depending on the types involved.

3.3 The Same Operator, Different Behavior

You noticed this at the end of Chapter 2, when `"2" + "3"` gave `"23"` instead of 5. Back then it was just an observation; now we can name the mechanism: the same symbol represents completely different operations, and Python decides which one by looking at the types of the values on either side. With numbers, `+` adds: `5 + 3` gives 8; with strings, it concatenates: `"5" + "3"` gives `"53"`.

This extends to `*`: with two numbers, it multiplies: `4 * 3` gives 12; with a string and an integer, it repeats: `"ha" * 3` gives `"hahaha"`. But with two strings? Python doesn't know what `"ha" * "3"` should mean, so it raises `TypeError`.




Figure 3.1: Same Operator, Different Behavior

When you mix integers and floats, Python promotes the result to a float: `5 + 3.0` gives `8.0`, not 8. This makes sense - a float can hold fractional values an integer can't, so Python picks the more general type and nothing gets lost.

The rest of the arithmetic from Chapter 2 - subtraction, division, `**`, `%`, and the precedence rules - works exactly as you saw it there, just with these type rules layered on top. One operator we skipped is worth adding now: `//` is integer division, which divides and throws away the remainder, so `10 // 3` gives 3 while `10 % 3` gives the leftover 1 - together they split a division into its whole-number part and its remainder.

That covers how operators respond to types. But floats themselves are hiding one more surprise - and since the receipt is about to do money math with them, it's worth seeing it now.

3.4 Why 0.1 + 0.2 Isn't 0.3

Try this in the REPL:

```
>>> 0.1 + 0.2
0.30000000000000004
```

That trailing 4 isn't a Python bug, and it isn't random - every mainstream programming language gives the same answer. To see why, we have to peek one level below the abstraction, at how the computer actually stores numbers.

Everything in a computer ultimately lives in binary. Remember the abstraction ladder from Chapter 1 that ended at electrical signals running through the processor? Those signals have two states, so every number the machine holds is written using only the digits 0 and 1. Whole numbers survive the translation perfectly - 6 becomes 110, 42 becomes 101010, nothing lost. Fractions are where it gets interesting.

Some fractions can't be written exactly in a given base, and you already know one: $1/3$ in our everyday base 10 is 0.3333..., the threes never end, and wherever you stop writing you've captured slightly less than $1/3$. Base 2 has the same problem, just with different fractions - and one tenth is one of them.

To see why, look at what binary digits after the point even are. In base 10 the places mean $1/10$, $1/100$, $1/1000$; in base 2 they mean $1/2$, $1/4$, $1/8$, $1/16$ - each place worth half the one before. Writing 0.1 in binary means building one tenth as a sum of these pieces. $1/2$, $1/4$, and $1/8$ are all too big; $1/16$ fits, add $1/32$ and you're at 0.09375; the next pieces that fit are $1/256$ and $1/512$, bringing you to 0.099609375 - closer, but not there. And it never gets there:




Figure 3.2: Building 0.1 in Binary

Written out as digits, that endless sum is 0.0001100110011..., the block 0011 repeating forever. A float has a fixed size - 64 bits in Python - so the summing has to stop, and what's stored is the closest value those pieces can build (the last bit gets rounded up rather than chopped, as it happens, which is why the stored value lands a hair above 0.1 instead of below). Those 64 bits are a memory budget: precision is being bought with space, and the error is what the budget can't cover. The same happens to 0.2. Add them, and the two tiny surpluses combine into something visibly different from 0.3; Python prints exactly what it's holding (the full storage format has a bit more machinery - a sign, an exponent - but you don't need it; the never-ending sum is the whole story).

This also tells you which decimals are safe: anything that is a finite sum of these pieces stores exactly. 0.5 is a single piece ($1/2$), 0.25 is another ($1/4$), 0.75 is two of them ($1/2 + 1/4$) - all perfect. 0.1, 0.2, and 0.3 never are.

For most purposes the error is far too small to matter, and later in this chapter : . 2 f will round it away before anyone sees it - the customer's receipt says \$0.30. For financial systems where every cent must be exact, Python has a `decimal` module that stores base-10 digits directly; we won't need it in this book, but know it exists.

When the error does matter and nobody accounts for it, the consequences can be real. In 1991, during the Gulf War, an American Patriot missile battery in Dhahran, Saudi Arabia failed to track and intercept an incoming Iraqi Scud, which struck an army barracks and killed 28 soldiers. The investigation traced the failure to exactly the number on this page: the system's clock counted time as a whole number of tenths of a second - and whole numbers, as we just saw, store perfectly. But to use that count in tracking calculations, the software multiplied it by 0.1 to convert to seconds, and the system's 24-bit hardware could only hold a chopped-off version of 0.1's endless binary sum, a constant off by about 0.000000095. Multiplied by the clock count after 100 hours of continuous operation, the computed time had drifted by about a third of a second, and a Scud covers more than half a kilometer in that time: the radar looked for the missile in the wrong place, found nothing, and dismissed the alert. The Army already knew about the drift - a corrected version of the software arrived in Dhahran the day after the attack.




Figure 3.3: The Patriot's Clock Drift

That's the one leak in the float abstraction worth knowing about. It still doesn't solve the receipt problem, though - the crash wasn't about precision, it was about strings pretending to be numbers. To fix that, we need to turn strings into numbers.

3.5 Type Conversion

Python provides built-in functions that convert values from one type to another.

`int()` converts to an integer. It works on strings that contain whole numbers, and on floats, but with a catch worth pausing on: it *truncates*, it doesn't round. It chops off everything after the decimal point and keeps the whole-number part, so `int(3.99)` is 3, not 4. That's a real trap - use `int()` on a price expecting it to round to the nearest value and you'll silently lose the fractional part. When rounding is what you actually want, that's a different function, `round()`: `round(3.99)` gives 4, and `round(3.14159, 2)` gives 3.14. Keep the two straight - `int()` chops, `round()` rounds.

```
>>> int("42")
42
>>> int(3.99)
3
>>> round(3.99)
4
```

`float()` converts to a float. It works on strings that contain numbers (with or without decimal points) and on integers:

```
>>> float("14.99")
14.99
>>> float(42)
42.0
```

`str()` converts anything to a string:

```
>>> str(42)
'42'
>>> str(3.14)
'3.14'
```

Now we can fix the receipt. If `quantity` is the string `"2"` and `price` is the string `"14.99"`, then:

```
>>> int("2") * float("14.99")
29.98
```

That's the total, and it's exactly the `float(price) * int(quantity)` we promised at the end of Chapter 2: `int()` turns the string `"2"` into the integer `2`, `float()` turns `"14.99"` into the float `14.99`, and the multiplication works because both values are now numbers.

You can convert at the moment you get the input:

```
quantity = int(input("Quantity: "))
price = float(input("Price per item: "))
```

Or convert later when you need the values as numbers - either way, the conversion has to happen before you do math.

What happens if the user types something that can't be converted? Try `int("hello")` in the REPL. Python raises a `ValueError`, and now you know how to read the traceback it prints. We'll learn how to handle these errors gracefully in a later chapter; for now, we'll trust that the user enters valid numbers.

Type conversion solves the arithmetic problem, but the receipt also needs to display values in a specific format - dollar signs, two decimal places, aligned columns - and string concatenation with `+` gets clumsy fast. There's a better way as we'll see.

3.6 Formatting Output with F-Strings

An **f-string** (formatted string literal) lets you embed expressions directly inside a string. Prefix the string with `f` and put expressions inside curly braces:

```
name = "Hannah"
age = 30
print(f"Hello, my name is {name} and I am {age} years old.")
```

Hello, my name is Hannah and I am 30 years old.

Compare that to the concatenation approach: "Hello, my name is " + name + " and I am " + str(age) + " years old." The f-string is shorter, more readable, and doesn't require converting age to a string manually.

You can put any expression inside the braces, not just variable names:

```
x = 10
y = 20
print(f"The sum of {x} and {y} is {x + y}")
```

The sum of 10 and 20 is 30

F-strings also support format specifiers after a colon, and this is where they become essential for the receipt.

To control decimal places, use `:.Nf` where N is the number of digits after the decimal point:

```
price = 14.99
tax = price * 0.08
print(f"Tax: ${tax:.2f}")
```

Tax: \$1.20

Without `:.2f`, the tax would display as `1.1992`, all four decimal places - and in larger calculations, the rounding error from earlier can surface as a long trail of digits. The format specifier rounds and fixes the display to exactly two decimal places, which is what you want for currency.

To add comma separators for large numbers, use `,:`:

```
population = 8000000
print(f"Population: {population:,}")
```

Population: 8,000,000

You can combine them: `${amount:,.2f}` gives you commas and two decimal places, which is exactly what currency formatting needs.

For alignment, use `<` (left), `>` (right), or `^` (center) with a width:

```
print(f"{'Item':<20}{'Price':>10}")
print(f"{'Book':<20}{'$14.99':>10}")
print(f"{'Notebook':<20}{'$3.50':>10}")
```

Item	Price
Book	\$14.99
Notebook	\$3.50

This aligns item names to the left in a 20-character field and prices to the right in a 10-character field - that's how receipts get their clean columns. If you did the badge exercise in Chapter 2, you've seen centering before: `^` does the same job as the `.center()` method, but inside an f-string, where it can sit alongside other formatting.

Now we have types, conversions, and formatting - almost enough to rebuild the receipt. But the receipt takes text input from the user, and user input is messy: someone might type " Bookshop " with extra spaces, or "hannah" in lowercase when you want "Hannah" on the receipt. We need tools to clean up and reshape strings before displaying them.

3.7 Working with Strings

Strings have methods: built-in features or tools (actually functions like the ones for converting data types, print, input; we'll talk more about this later) you call with a dot after the string. You don't need to memorize all of them, but a handful show up constantly.

Case conversion: `"hello".upper()` gives "HELLO", `"HELLO".lower()` gives "hello", `"hello world".title()` gives "Hello World". We'll use `.title()` in the receipt to normalize names regardless of how the user types them, and `.strip()` to remove accidental spaces.

Checking content: `"Hello".startswith("He")` returns `True`, `"Hello".endswith("lo")` returns `True`, `"fun" in "Python is fun"` returns `True`. These let you inspect strings without manually picking through characters.

Finding and replacing: `"Hello, world!".find("world")` returns 7 (the position where "world" starts), and `"Hello, world!".replace("world", "Python")` returns "Hello, Python!". Note that `replace()` doesn't change the original string - it returns a new one. Strings in Python are **immutable**: once created, they can't be modified, and every operation that appears to change a string actually creates a new one.

Splitting and joining: `"apple,banana,cherry".split(",")` returns `['apple', 'banana', 'cherry']` (a list, which we'll cover properly in a later chapter), and `",".join(['apple', 'banana', 'cherry'])` goes the other way: `"apple, banana, cherry"`. The split/join pair is how you take structured text apart and put it back together.

Stripping whitespace: `" hello ".strip()` returns "hello", removing spaces (and tabs, new-lines) from both ends; `lstrip()` and `rstrip()` remove from the left or right only. This is essential when processing user input, which often has accidental spaces.

Length: `len("hello")` returns 5. You met this in the Chapter 2 postcard exercise - it's a built-in function, not a method, so you call it with the string as an argument rather than using dot notation.

Individual characters: `"hello"[0]` gives `"h"`, `"hello"[-1]` gives `"o"`. You used `name[0]` for the initials exercise in Chapter 2; now the full picture: Python uses zero-based indexing (the first character is at position 0, the second at position 1) and negative indices count from the end.

Slicing: `"hello"[1:4]` gives `"ell"` (characters at positions 1, 2, and 3 - the end index is excluded), `"hello"[:3]` gives `"hel"`, `"hello"[3:]` gives `"lo"`. This lets you extract portions of a string by position.




Figure 3.4: Indexing and Slicing

That's enough string tools to build something real. Let's go back to the receipt.

3.8 Hands-On: Receipt Printer, Completed

The receipt from Chapter 2 collected input and formatted it as text, but couldn't do any math, now we can fix that. We'll rebuild it with type conversions for arithmetic, string methods to clean up input, and f-strings for formatting. If you did the first exercise in Chapter 2, you already extended the receipt to three line items - that's exactly the skeleton we're rebuilding here, this time with working math.

Start a fresh project (keep your Chapter 2 folder as it is - we'll want to compare against it later):

```
uv init chapter3
cd chapter3
```

Create a file called `receipt.py`. Start with the input, converting numbers at the point of entry and cleaning up text with string methods. We're also moving the customer name up next to the store name: with three items coming, it's cleaner to collect all the header information first:


```
# receipt.py
print("=== Receipt Printer ===", end="\n\n")

store_name = input("Store name: ").strip().title()
customer_name = input("Customer name: ").strip().title()
print()
```

Then the three items:

```
# receipt.py
# ...existing code above

item1_name = input("Item 1 name: ").strip()
item1_qty = int(input("Item 1 quantity: "))
item1_price = float(input("Item 1 price: "))

item2_name = input("Item 2 name: ").strip()
item2_qty = int(input("Item 2 quantity: "))
item2_price = float(input("Item 2 price: "))

item3_name = input("Item 3 name: ").strip()
item3_qty = int(input("Item 3 quantity: "))
item3_price = float(input("Item 3 price: "))
```

Notice what changed from Chapter 2. Quantities go through `int()` and prices through `float()` - they're numbers now, not strings, which means we can compute line totals and a grand total. The store and customer names get `.strip()` (removing accidental spaces) and `.title()` (capitalizing properly), so even if the user types " bookshop ", the receipt prints "Bookshop." Item names get `.strip()` too.

Now the calculations:

```
# receipt.py
# ...existing code above

item1_total = item1_qty * item1_price
item2_total = item2_qty * item2_price
item3_total = item3_qty * item3_price

subtotal = item1_total + item2_total + item3_total
tax_rate = 0.08
tax = subtotal * tax_rate
total = subtotal + tax
```

Just six lines of arithmetic - in Chapter 2 this was impossible, now it works because the values are the right types.

Finally, the formatted output, we'll use f-strings with alignment to make clean columns. To get the dollar sign and the number to align as a unit, we format each price as a string first, then right-align that string in a fixed-width field:

```
# receipt.py
# ...existing code above

width = 35

print()
print("=" * width)
print(f"{{store_name:^^{{width}}}}")
print("=" * width)
print(f"Customer: {{customer_name}}")
print("-" * width)

print(f"{{'Item':<15}} {{'Qty':>3}} {{'Price':>7}} {{'Total':>7}}")
print("-" * width)
```

Now the three item rows - format each price with its dollar sign first, then right-align the result:

```
# receipt.py
# ...existing code above

p1 = f"${item1_price:.2f}"
t1 = f"${item1_total:.2f}"
print(f"{{item1_name:<15}} {{item1_qty:>3}} {{p1:>7}} {{t1:>7}}")

p2 = f"${item2_price:.2f}"
t2 = f"${item2_total:.2f}"
print(f"{{item2_name:<15}} {{item2_qty:>3}} {{p2:>7}} {{t2:>7}}")

p3 = f"${item3_price:.2f}"
t3 = f"${item3_total:.2f}"
print(f"{{item3_name:<15}} {{item3_qty:>3}} {{p3:>7}} {{t3:>7}}")
```

And the totals section, closing the frame:

```
# receipt.py
# ...existing code above

print("-" * width)
sub_str = f"${subtotal:.2f}"
tax_str = f"${tax:.2f}"
total_str = f"${total:.2f}"
print(f"{{'Subtotal:':>27}} {{sub_str:>7}}")
print(f"{{'Tax (8%):':>27}} {{tax_str:>7}}")
print("=" * width)
print(f"{{'TOTAL:':>27}} {{total_str:>7}}")
print("=" * width)
print(f"{{'Thank you for your purchase!':^{{width}}}}")
```

Run it with `uv run receipt.py` (or `uv run python receipt.py` - both work; `uv run` will hand a `.py` file straight to Python either way). Enter “bookshop” as the store and “hannah” as the customer (both lowercase, to see `.title()` at work), then three items: “Python Book” at quantity 2, price 29.99; “Notebook” at quantity 3, price 4.50; “Pen” at quantity 5, price 1.25. The output looks like:

```

=====
                        Bookshop
=====
Customer: Hannah
-----
Item                Qty    Price    Total
-----
Python Book         2    $29.99    $59.98
Notebook             3     $4.50    $13.50
Pen                  5     $1.25     $6.25
-----
                        Subtotal:   $79.73
                        Tax  (8%):   $6.38
=====
                        TOTAL:    $86.11
=====
Thank you for your purchase!

```

Compare this to the Chapter 2 version: that receipt was a string-formatting exercise - it could display prices but not add them; this receipt computes real totals, applies tax, cleans up input, and formats everything with aligned columns and two decimal places. The difference is type conversion (turning input strings into numbers), string methods (cleaning and normalizing text), and f-strings (controlling how values display).

The rounding error from earlier is quietly at work here too: the stored subtotal is actually `79.72999999999999`, and the total is `86.10839999999999`. But `:.2f` rounds for display, so the customer sees `$79.73` and `$86.11`.

Look at how the columns work: each price gets formatted as a string with a dollar sign first (`f"${item1_price:.2f}"` produces `"$29.99"`), then that string gets right-aligned in a 7-character field (`{p1:>7}` produces `" $29.99"` or `" $4.50"`); since the header uses the same field widths (`{'Price':>7}`), everything lines up. The two-step approach - format the value, then align the formatted string - is a pattern you'll use whenever you need columns with mixed content like currency symbols.

Try experimenting: change the tax rate, add a discount line. What happens if an item name is longer than 15 characters? What happens if a price is over \$999.99 (it would need more than 7 characters)? The format specifiers control the layout, and testing their limits teaches you what they can and can't handle.

The receipt works, but there's an obvious problem: we hardcoded three items. What if the customer buys two items? What if they buy ten? Right now you'd need a separate set of variables for each item, and you'd need to know the count before writing the code. In Chapter 4, we'll learn how to make decisions (what if the quantity is zero?) and in Chapter 5, how to repeat actions (process as many items as the customer has).

3.9 Chapter Summary

The receipt crashed in Chapter 2 because `input()` returns strings, and strings can't do arithmetic; the traceback told us exactly that, once we learned to read it. The fix was understanding types: Python distinguishes between integers, floats, strings, and booleans, and the type of a value determines what operations it supports. The same operator behaves differently depending on the types involved - `+` adds numbers but concatenates strings, `*` multiplies numbers but repeats strings. And floats are binary approximations: 0.1 has no exact binary form, so `0.1 + 0.2` carries an error in the sixteenth decimal place, but it's mostly harmless once display rounding is in place.

Type conversion helps us: `int()`, `float()`, and `str()` let you move values between types when you need to: `int(input("Quantity: "))` turns the user's text into a number you can compute with. F-strings handle the other direction, embedding numbers inside formatted text with control over decimal places, alignment, and width.

Strings are immutable sequences of characters with built-in methods for case conversion, searching, replacing, splitting, and stripping whitespace. You access individual characters by index and extract substrings by slicing - tools that show up constantly once you start processing text.

The rebuilt receipt computes real totals with tax and formatted columns, and cleans up user input with `.strip()` and `.title()`, but it handles exactly three items, hardcoded. What if the customer has more? What if they have fewer? The program can't adapt because it follows the same path every time. In the next chapter, we'll introduce conditionals: the ability to make decisions based on data, so your programs can respond to different situations differently.

3.10 Exercises

1. The receipt uses a fixed tax rate of 8%. Modify the program to ask the user for the tax rate as a percentage (like 8.5), convert it to a decimal, and apply it. Then add a second version of the total that shows what the bill would be with no tax. Print both totals side by side. How does the formatting need to change when the tax rate has a decimal place?
2. Write a program that asks for a distance in kilometers and converts it to miles, feet, and inches (1 km = 0.621371 miles, 1 mile = 5280 feet, 1 foot = 12 inches). Display all four values aligned in a column with appropriate precision: kilometers and miles to two decimal places, feet to one, inches to zero. What happens to the inches value when you enter a very small distance like 0.001 km?
3. Write a program that asks for a full name (like "jane doe") and prints it in four formats: all uppercase, all lowercase, title case, and as initials with periods (like "J.D."). You'll need `split()` to separate the words and indexing (`word[0]`) to get first characters. What happens if the user enters a name with three parts, like "Mary Jane Watson"? What about a single name with no space?
4. Write a program that asks the user for a sentence and reports: the number of characters, the number of words (use `split()`), the first word, the last word, and the sentence reversed (use slicing with

[: : -1]). The reversed sentence will look like gibberish, but compare it to the original: is a palindrome sentence (one that reads the same forwards and backwards) possible with whole words? Try entering “was it a car or a cat i saw” and check.

5. F-strings let you create formatted tables. Write a program that asks for three students’ names and their scores on two tests. Display the results as a table with right-aligned scores, each student’s average to one decimal place, and a class average at the bottom. The tricky part is getting the columns to line up. What’s the minimum field width that keeps the table readable for any reasonable name length?
6. Write a program that asks for a price in dollars (like 1234 . 56) and prints it formatted three ways: plain (1234 . 56), with commas (1 , 234 . 56), and as a full currency string (\$1 , 234 . 56). Then ask for a second price and display both prices right-aligned in a 15-character field, one above the other, with a line of dashes between them and a total below. This mimics the column-alignment logic from the receipt but requires you to figure out the format specifiers yourself.

4 Making Decisions

The receipt from Chapter 3 computes totals, formats columns, and cleans up input, but it follows the same path every time it runs: ask for three items, calculate, print. Type a quantity of zero and the item still lands on the bill; buy \$200 worth of books and no discount appears, because the program has no way to notice any of it. Every run walks the same road.

What's missing is the ability to make decisions: examine a condition and take different actions depending on whether it's true or false, and that's what this chapter introduces. By the end, you'll rebuild the receipt so it validates what the user types, skips items that shouldn't appear, applies tiered discounts, and shapes its own output to match - the first version that behaves differently depending on what happens (the receipt's other rigidity - exactly three items, no more, no fewer - is a repetition problem instead of a decision problem; that one falls in Chapter 5).

4.1 True, False, and Comparisons

Before we can make decisions in code, we need a way to express conditions. Chapter 3 introduced the type for this: the boolean, with its two values `True` and `False`.

```
>>> type(True)
<class 'bool'>
>>> type(False)
<class 'bool'>
```

You rarely type `True` or `False` directly - usually, booleans come from comparisons. Python provides six comparison operators:

```
>>> 5 == 5
True
>>> 5 == 3
False
>>> 5 != 3
True
>>> 5 > 3
True
>>> 5 < 3
False
>>> 5 >= 5
True
>>> 5 <= 4
False
```

`==` tests equality, and a common beginner mistake is confusing it with `=`, which is assignment: `x = 5` assigns the value 5 to `x` (you can loosely think of this as storing the value 5 in `x`, but we'll see later why that's not exactly correct), while `x == 5` asks whether `x` is currently 5 and returns `True` or `False`.

Comparisons work on strings too: `"alice" == 'alice'` is `True`, `"alice" == "Alice"` is `False`, and `"apple" < "banana"` is `True` - alphabetical order, the same ordering Chapter 1's phone book was sorted by. Case sensitivity is why we sometimes normalize inputs to compare them: the receipt already uses `.strip()` and `.title()` to clean what it displays, and `.lower()` is used when a program needs to check what the user meant regardless of whether they typed upper or lower - this chapter's exercises will lean on exactly that.

These comparisons produce the booleans that determine the branch decisions, so the question is what to do with them.

4.2 The if Statement

The simplest decision: do something only if a condition is true.

```
temperature = float(input("Temperature: "))

if temperature > 30:
    print("It's hot today.")
    print("Stay hydrated.")

print("Have a good day.")
```

If the user enters 35, both messages print; if they enter 20, only "Have a good day" prints (this is in Celsius, of course). The indented block after `if` runs only when the condition is `True`, and the unindented line after runs regardless. So with temperature 35, Python sees:

```
if 35 > 30:
    ...
```

which is `True`, so it's equivalent to:

```
if True:
    ...
```

and it executes the code inside.

The indentation is how Python knows which lines belong to the `if`. Every line indented under it (by convention, 4 spaces) is part of the block, and the first line that returns to the previous indentation level is outside it. This is different from many other languages, which use braces `{ }` to mark blocks - in Python, the whitespace is the syntax.

```
if temperature > 30:
    print("This is inside the if block.")
    print("So is this.")
print("This is outside. It always runs.")
```

If you forget to indent, Python raises an `IndentationError` - a traceback like any other, and by now you know to read its last line. Inconsistent indentation (4 spaces here, 3 there) gets the same error. Use 4 spaces per level and let your editor handle it.

Getting an `IndentationError` is the best error to get though: you fix the indentation and you're good to go. The other scenario is your code running and the logic not working, which could happen if you had something like:

```
temperature = float(input("Temperature: "))

if temperature > 30:
    print("It's hot today.")
print("Stay hydrated.")

print("Have a good day.")
```

This will always tell people to stay hydrated instead of only when it's hot (not a problem here, but it could be in another program) - an issue with the logic instead of the syntax, Chapter 1's algorithm-versus-implementation split showing up in miniature.

An `if` answers the question "should I do this?" But sometimes you need to choose between two alternatives.

4.3 Choosing Between Two Paths

The `if-else` statement handles "do this or do something else":

```
age = int(input("Your age: "))

if age >= 18:
    print("You can vote.")
else:
    print("You can't vote yet.")
```

Exactly one of the two blocks runs: the `if` block when the condition is `True`, the `else` block when it's `False`. There's no scenario where both run or neither runs.

This is the pattern for the receipt's input validation - when the user types a price, we want to check whether it makes sense before using it. We can't do that fully yet (a user typing "abc" still crashes `float()` - handling that cleanly needs `try/except`, which comes in a later chapter, though you could laboriously pre-check the text with the string methods you already have), but we can check the conditions we do understand:


```
price = float(input("Price: "))

if price < 0:
    print("Price can't be negative.")
else:
    print(f"Price accepted: ${price:.2f}")
```

This gives us some validation already, but what about situations with more than two possibilities?

4.4 Multiple Conditions with elif

The grading scale is a classic example: a score maps to one of several grades; you could write this with nested if-else statements:

```
score = int(input("Score: "))

if score >= 90:
    print("Grade: A")
else:
    if score >= 80:
        print("Grade: B")
    else:
        if score >= 70:
            print("Grade: C")
        else:
            if score >= 60:
                print("Grade: D")
            else:
                print("Grade: F")
```

But Python provides a cleaner `elif` (short for “else if”) for chains of conditions:

```
score = int(input("Score: "))

if score >= 90:
    print("Grade: A")
elif score >= 80:
    print("Grade: B")
elif score >= 70:
    print("Grade: C")
elif score >= 60:
    print("Grade: D")
else:
    print("Grade: F")
```

Python checks each condition in order, and the moment it finds one that’s `True`, it runs that block and skips the rest. If none are `True`, the `else` block runs.

The order matters - if you wrote `if score >= 60` first, everyone with 60 or above would get a D, because that condition would be `True` for scores of 90 too, and Python would stop there. The chain

works because it checks the most restrictive condition first:




Figure 4.1: The elif Chain: First True Wins

The `else` at the end is optional. You can have an `if-elif` chain with no `else`, in which case it's possible for nothing to run if all conditions are `False` - but when you're categorizing something exhaustively (every score gets a grade), the `else` acts as a catch-all.

Now let's think about the receipt: the tax rate might vary by location, and a store might offer different discount tiers based on the subtotal. An `if-elif-else` chain handles that:

```
if subtotal >= 100:
    discount = 0.10
elif subtotal >= 50:
    discount = 0.05
else:
    discount = 0.0
```

Single conditions work for one test at a time, but real decisions often depend on multiple things at once - you can drive if you're old enough *and* you have a license, you should bring an umbrella if it's raining

or the forecast says rain - and these require combining conditions.

4.5 Combining Conditions

Python has three logical operators for combining booleans: and, or, and not.

and is True only when both sides are True:

```
age = 25
has_license = True

if age >= 18 and has_license:
    print("You can drive.")
```

Both conditions must hold - if either is False, the whole expression is False.

Note that you can also add parentheses to the conditions to make them more readable:

```
age = 25
has_license = True

if (age >= 18) and has_license:
    print("You can drive.")
```

This is equivalent to the version with no parentheses.

or is True when at least one side is True:

```
is_weekend = True
is_holiday = False

if is_weekend or is_holiday:
    print("No work today.")
```

Only one needs to hold; the expression is False only when both sides are False.

not inverts a boolean:

```
is_raining = False

if not is_raining:
    print("Leave the umbrella at home.")
```

You can combine these freely, but when mixing and and or, and has higher precedence (it binds tighter):

```
if user_type == "admin" or user_type == "editor" and is_approved:
    grant_access()
```

Since `and` binds tighter, Python reads this as `user_type == "admin" or (user_type == "editor" and is_approved)` - admins get access regardless of approval. If you meant to require approval for both, you need parentheses:

```
if (user_type == "admin" or user_type == "editor") and is_approved:
    grant_access()
```

When in doubt, use parentheses; they make your intent explicit and cost nothing.

One useful behavior: Python evaluates `and` and `or` lazily. In `x and y`, if `x` is `False`, Python doesn't bother evaluating `y` (the result is `False` either way), and in `x or y`, if `x` is `True`, Python skips `y`. This is called **short-circuit evaluation**, and it lets you write safe checks like:

```
if age and age > 0:
    print("Age is positive.")
```

If `age` is `None` (say the user skipped the question), it's falsy - treated as `False` - so the `and` short-circuits and Python never evaluates `age > 0`. That's the crash that never runs: comparing `None` to a number raises a `TypeError`, but the right side is never touched. You don't need this pattern yet, but you'll see it often in other people's code.




Figure 4.2: Short-Circuit Evaluation: Stopping Early

Python considers several values as `False` in boolean contexts - `0`, `0.0`, `""` (the empty string), `None`,

and empty collections; everything else is truthy. This means `if name:` is a common shorthand for “if name is not empty.”

For the receipt, combining conditions lets us check multiple things at once: is the quantity positive *and* is the price valid? Is the subtotal above the discount threshold *or* does the customer have a loyalty card? These compound conditions help you make nuanced decisions.

We’ve been writing flat `if-elif-else` chains so far, but sometimes nested conditionals are the best approach.

4.6 Nested Conditionals

As we saw with the grading scale, you can put an `if` inside another `if`:

```
weather = input("Weather (rainy/sunny): ").lower()
temp = float(input("Temperature: "))

if weather == "rainy":
    if temp < 10:
        print("Heavy coat and umbrella.")
    else:
        print("Light jacket and umbrella.")
else:
    if temp > 30:
        print("Sunscreen and water.")
    else:
        print("Enjoy the day.")
```

Each level of nesting adds another 4 spaces of indentation. This works, but deep nesting gets hard to read quickly: two levels are common and fine, three are a warning sign, four or more usually means you should restructure (notice the `.lower()` on the first line, doing exactly the normalization job from earlier - “Rainy”, “RAINY”, and “rainy” all become the same string before any comparison happens).

Often you can flatten nested conditions by combining them with `and`:

```
# Nested
if weather == "rainy":
    if temp < 10:
        print("Heavy coat and umbrella.")

# Flat
if weather == "rainy" and temp < 10:
    print("Heavy coat and umbrella.")
```

The flat version is better when there’s only one action; the nested version is better when the outer condition branches into multiple inner decisions, like the weather example above, where “rainy” leads to different temperature checks than “sunny” does.

That’s enough to build something real, let’s put conditionals to work on the receipt.

4.7 Hands-On: Receipt Printer with Validation and Discounts

The Chapter 3 receipt computed totals but had no ability to adapt, now we'll add three capabilities: validating that prices and quantities make sense, skipping items with zero quantity, and applying a discount based on the subtotal.

Start a fresh project:

```
uv init chapter4
cd chapter4
```

Create `receipt.py`. We can't re-ask the user when they enter bad data yet, but we can check for obviously wrong values and refuse to use them. Start with the header:

```
# receipt.py
print("=== Receipt Printer ===", end="\n\n")

store_name = input("Store name: ").strip().title()
customer_name = input("Customer name: ").strip().title()
```

Now the items; we'll collect three, but check each one: if a quantity is zero, we mark the item to skip later; if anything is negative, we reject it. Here's the first:

```
# receipt.py
# ...existing code above

print()
item1_name = input("Item 1 name: ").strip()
item1_qty = int(input("Item 1 quantity: "))
item1_price = float(input("Item 1 price: "))
item1_valid = True

if item1_qty < 0 or item1_price < 0:
    print(" Invalid: quantity and price must not be negative.")
    item1_valid = False
elif item1_qty == 0:
    print(f" Skipping {item1_name} (quantity is 0).")
    item1_valid = False
```

The `if-elif` chain does the validation: negative values are rejected outright, zero quantity is accepted but skipped, and the `_valid` flag records the verdict for the rest of the program to consult. Notice `or` combining two checks in the first condition - either a negative quantity or a negative price makes the item invalid.

Items 2 and 3 are the exact same block with the numbers changed:

```
# receipt.py
# ...existing code above
```

```

print()
item2_name = input("Item 2 name: ").strip()
item2_qty = int(input("Item 2 quantity: "))
item2_price = float(input("Item 2 price: "))
item2_valid = True

if item2_qty < 0 or item2_price < 0:
    print(" Invalid: quantity and price must not be negative.")
    item2_valid = False
elif item2_qty == 0:
    print(f" Skipping {item2_name} (quantity is 0).")
    item2_valid = False

```

And once more for item 3:

```

# receipt.py
# ...existing code above

print()
item3_name = input("Item 3 name: ").strip()
item3_qty = int(input("Item 3 quantity: "))
item3_price = float(input("Item 3 price: "))
item3_valid = True

if item3_qty < 0 or item3_price < 0:
    print(" Invalid: quantity and price must not be negative.")
    item3_valid = False
elif item3_qty == 0:
    print(f" Skipping {item3_name} (quantity is 0).")
    item3_valid = False

```

Typing that block twice more felt mechanical for a reason - hold that thought until the end of the chapter.

Now the calculations, including only valid items in the subtotal:

```

# receipt.py
# ...existing code above

item1_total = item1_qty * item1_price if item1_valid else 0
item2_total = item2_qty * item2_price if item2_valid else 0
item3_total = item3_qty * item3_price if item3_valid else 0

subtotal = item1_total + item2_total + item3_total

```

The `x if condition else y` pattern is a **conditional expression** (sometimes called a ternary operator). It evaluates to `x` when the condition is `True` and `y` when it's `False` - here, it computes the line total only for valid items and uses 0 otherwise. This is useful for one-line decisions; for anything more complex, stick with a regular `if-else` block.

Now the discount, tiered by subtotal:

```
# receipt.py
# ...existing code above

if subtotal >= 100:
    discount_rate = 0.10
    discount_label = "10%"
elif subtotal >= 50:
    discount_rate = 0.05
    discount_label = "5%"
else:
    discount_rate = 0.0
    discount_label = ""
```

The if-elif-else chain picks the tier: \$100 or more gets 10%, \$50 or more gets 5%, under \$50 gets nothing - first true condition wins, exactly like the grading scale. The rest of the math follows from the chosen rate:

```
# receipt.py
# ...existing code above

discount = subtotal * discount_rate
after_discount = subtotal - discount

tax_rate = 0.08
tax = after_discount * tax_rate
total = after_discount + tax
```

Finally, the formatted output. Conditionals control what appears: only valid items get printed, and the discount lines only show when there is one. Start with the frame - the same header and column layout as Chapter 3:

```
# receipt.py
# ...existing code above

width = 35

print()
print("=" * width)
print(f"{{store_name:^{width}}}")
print("=" * width)
print(f"Customer: {{customer_name}}")
print("-" * width)
print(f"{{'Item':<15}} {{'Qty':>3}} {{'Price':>7}} {{'Total':>7}}")
print("-" * width)
```

Now the item rows. Each one sits inside an if with no else - invalid items don't print an apology, they simply never appear:

```
# receipt.py
# ...existing code above
```



```

if item1_valid:
    p1 = f"${item1_price:.2f}"
    t1 = f"${item1_total:.2f}"
    print(f"{item1_name:<15} {item1_qty:>3} {p1:>7} {t1:>7}")

```

Rows 2 and 3 are the same guard with the numbers changed:

```

# receipt.py
# ...existing code above

if item2_valid:
    p2 = f"${item2_price:.2f}"
    t2 = f"${item2_total:.2f}"
    print(f"{item2_name:<15} {item2_qty:>3} {p2:>7} {t2:>7}")

if item3_valid:
    p3 = f"${item3_price:.2f}"
    t3 = f"${item3_total:.2f}"
    print(f"{item3_name:<15} {item3_qty:>3} {p3:>7} {t3:>7}")

```

Below the rows, the subtotal always prints, but the entire discount section exists only when a discount fired:

```

# receipt.py
# ...existing code above

print("-" * width)
sub_str = f"${subtotal:.2f}"
print(f"{'Subtotal:':>27} {sub_str:>7}")

if discount_rate > 0:
    disc_str = f"${discount:.2f}"
    print(f"{'Discount (' + discount_label + ')':>27} {disc_str:>7}")
    after_str = f"${after_discount:.2f}"
    print(f"{'After discount:':>27} {after_str:>7}")

```

And the close is unconditional - every receipt gets tax, a total, and the thank-you:

```

# receipt.py
# ...existing code above

tax_str = f"${tax:.2f}"
total_str = f"${total:.2f}"
print(f"{'Tax (8%)':>27} {tax_str:>7}")
print("=" * width)
print(f"{'TOTAL:':>27} {total_str:>7}")
print("=" * width)
print(f"{'Thank you for your purchase!':^{width}}")

```

Run it with `uv run receipt.py`. Enter “bookshop” and “hannah” as before, then three items: “Book” at quantity 2, price 29.99; “Pen” at quantity 5, price 1.25; “Notebook” at quantity 3, price 4.50. The output should look something like:

```

=====
                        Bookshop
=====
Customer: Hannah
-----
Item                Qty   Price   Total
-----
Book                 2   $29.99   $59.98
Pen                  5    $1.25    $6.25
Notebook             3    $4.50   $13.50
-----
                        Subtotal:  $79.73
                        Discount (5%):  -$3.99
                        After discount:  $75.74
                        Tax (8%):    $6.06
=====
                        TOTAL:   $81.80
=====
Thank you for your purchase!

```

Same items as Chapter 3, but a different receipt now: the subtotal crossed \$50, so the 5% tier fired and two new lines appeared that the program decided to print. Now run it again and try the other paths; set one item's quantity to 0 - it vanishes from the receipt entirely, and the totals adjust. Try a negative price - the item is rejected with a message; buy items totaling under \$50 - the discount lines don't appear at all. One program, but from the data it creates four different receipts.

The conditionals do three different jobs here. The `if item_valid` checks control which items appear (conditional display), the `if-elif-else` chain picks the discount tier (conditional logic), and the `if discount_rate > 0` check controls whether the discount section exists at all (conditional formatting). Same mechanism, different uses.

Look at how much more capable this receipt is compared to Chapter 3's version: it validates input, adapts its display, and applies business logic. But there's still the repetition problem - the same block of code appears three times for three items, with only the variable names changing. That's tedious to write, error-prone to maintain, and impossible to scale: a customer with 20 items would need 20 copies of the same block, and you'd have to know the count before writing the code. In Chapter 5, we'll introduce loops, which let you write the block once and repeat it as many times as needed.

4.8 Chapter Summary

The receipt from Chapter 3 followed the same path every time, and conditionals give programs the ability to choose: examine a condition and take different actions depending on whether it's true or false.

Booleans - the two-valued type Chapter 3 planted - come from comparisons (`==`, `!=`, `>`, `<`, `>=`, `<=`) and combine with `and`, `or`, and `not`. When mixing those operators, `and` binds tighter than `or`, so use

parentheses to make your intent explicit. Python evaluates them lazily: short-circuit evaluation means the second half of an `and` or `or` might never run at all.

The `if` statement runs a block only when a condition is true, `if-else` chooses between two paths, and `if-elif-else` handles multiple conditions checked in order - the first true condition wins, and everything after it is never looked at. Chains can be nested, but flattening with `and` is often clearer. Indentation is Python's syntax for marking all of these blocks: 4 spaces per level, consistently, or the traceback will tell you otherwise.

The rebuilt receipt uses conditionals for three purposes: validating input (rejecting negatives, skipping zero quantities), applying business logic (tiered discounts), and controlling display (only valid items appear, and the discount lines exist only when earned). But the same code still appears three times for three items. Loops, which we'll introduce in Chapter 5, eliminate that repetition by letting you write a block once and execute it as many times as needed - and once the item count stops being fixed, the receipt starts becoming a real point of sale.

4.9 Exercises

1. The receipt applies a discount based on the subtotal; add a second discount mechanism: a loyalty code. Ask the user for a code before the items - if they enter "SAVE20", apply an additional 20% discount after the tier discount; if they enter "HALF", apply 50%; anything else (including empty) means no loyalty discount. Decide how to handle case - should "save20" work? - and use the normalization tools from Chapter 3 to make your decision real. Print a message confirming which code was applied (or that none was). How do the two discounts interact? Should the loyalty discount apply to the original subtotal or to the already-discounted amount?
2. Write a program that asks for three numbers and prints them in order from smallest to largest. You'll need to compare all three pairs and figure out the ordering; note that there are six possible orderings of three numbers. How many comparisons does your program make in the worst case? Can you do it in exactly three comparisons for any input?
3. Write a program that asks for a year and determines whether it's a leap year. The rule: a year is a leap year if it's divisible by 4, except years divisible by 100 are not leap years, unless they're also divisible by 400. So 2000 is a leap year, 1900 is not, and 2024 is. Express this as a single boolean expression using `and`, `or`, and `not`; then write it as an `if-elif-else` chain. Which version is easier to read? Test with 2000, 1900, 2024, and 2023.
4. Write a password strength checker. Ask for a password and rate it as "weak", "medium", or "strong" based on these criteria: at least 8 characters long, contains at least one digit (check each character with `.isdigit()`), and contains at least one uppercase letter (`.isupper()`). Weak means it fails two or more criteria, medium means it fails one, strong means it passes all three. You'll need to check each criterion separately, count how many pass, then use that count to determine the rating. What happens with an empty password?

5. Write a temperature converter that asks for a temperature and its unit (C or F), then converts to the other. If the input unit is "C", convert to Fahrenheit with $F = C \times 9/5 + 32$; if it's "F", convert to Celsius with $C = (F - 32) \times 5/9$; if the user enters anything else, print an error. Add a check: if the converted temperature is below absolute zero (-273.15°C or -459.67°F), warn the user that the input is physically impossible. What's the relationship between the two absolute zero values?
6. Write a calculator that asks for two numbers and an operation (+, -, *, /), then performs the operation and prints the result. Handle division by zero: if the user tries to divide by zero, print an error instead of crashing. Then add a second check: if the result is a whole number (like 6.0), display it without the decimal point (as 6). Use the conditional expression `int(result) if result == int(result) else result` and explain to yourself why that works. What happens with very large numbers?

Downey, Allen B. 2015. *Think Python: How to Think Like a Computer Scientist*. 2nd ed. O'Reilly Media.

Gutttag, John. 2016. *Introduction to Computation and Programming Using Python: With Application to Understanding Data*. 2nd ed. MIT Press.

Hammack, Richard H. 2020. *Book of Proof*. 3rd ed. Virginia Commonwealth University.

Percival, Harry J. W., and Bob Gregory. 2020. *Architecture Patterns with Python: Enabling Test-Driven Development, Domain-Driven Design, and Event-Driven Microservices*. O'Reilly Media.

Pólya, George. 2014. *How to Solve It: A New Aspect of Mathematical Method*. 2nd ed. Princeton University Press.

Ramalho, Luciano. 2021. *Fluent Python: Clear, Concise, and Effective Programming*. 2nd ed. O'Reilly Media.

Robson, Patrick. 2021. *Robust Python: Write Clean and Maintainable Code*. O'Reilly Media.

Shaw, Zed A. 2013. *Learn Python the Hard Way: A Very Simple Introduction to the Terrifyingly Beautiful World of Computers and Code*. 3rd ed. Zed Shaw's Hard Way Series. Addison-Wesley Professional.

Slatkin, Brett. 2019. *Effective Python: 125 Specific Ways to Write Better Python*. 2nd ed. Effective Software Development Series. Addison-Wesley Professional.

Sweigart, Al. 2019. *Automate the Boring Stuff with Python: Practical Programming for Total Beginners*. 2nd ed. No Starch Press.

Velleman, Daniel J. 2019. *How to Prove It: A Structured Approach*. 3rd ed. Cambridge University Press.